

Complex Computing Problem

OPTIMIZING A KLT FEATURE TRACKER USING GPU KERNELS

Menahil Ahmad 23I-0546
Hiyam Rehan 23i-0039

Experimental Setup

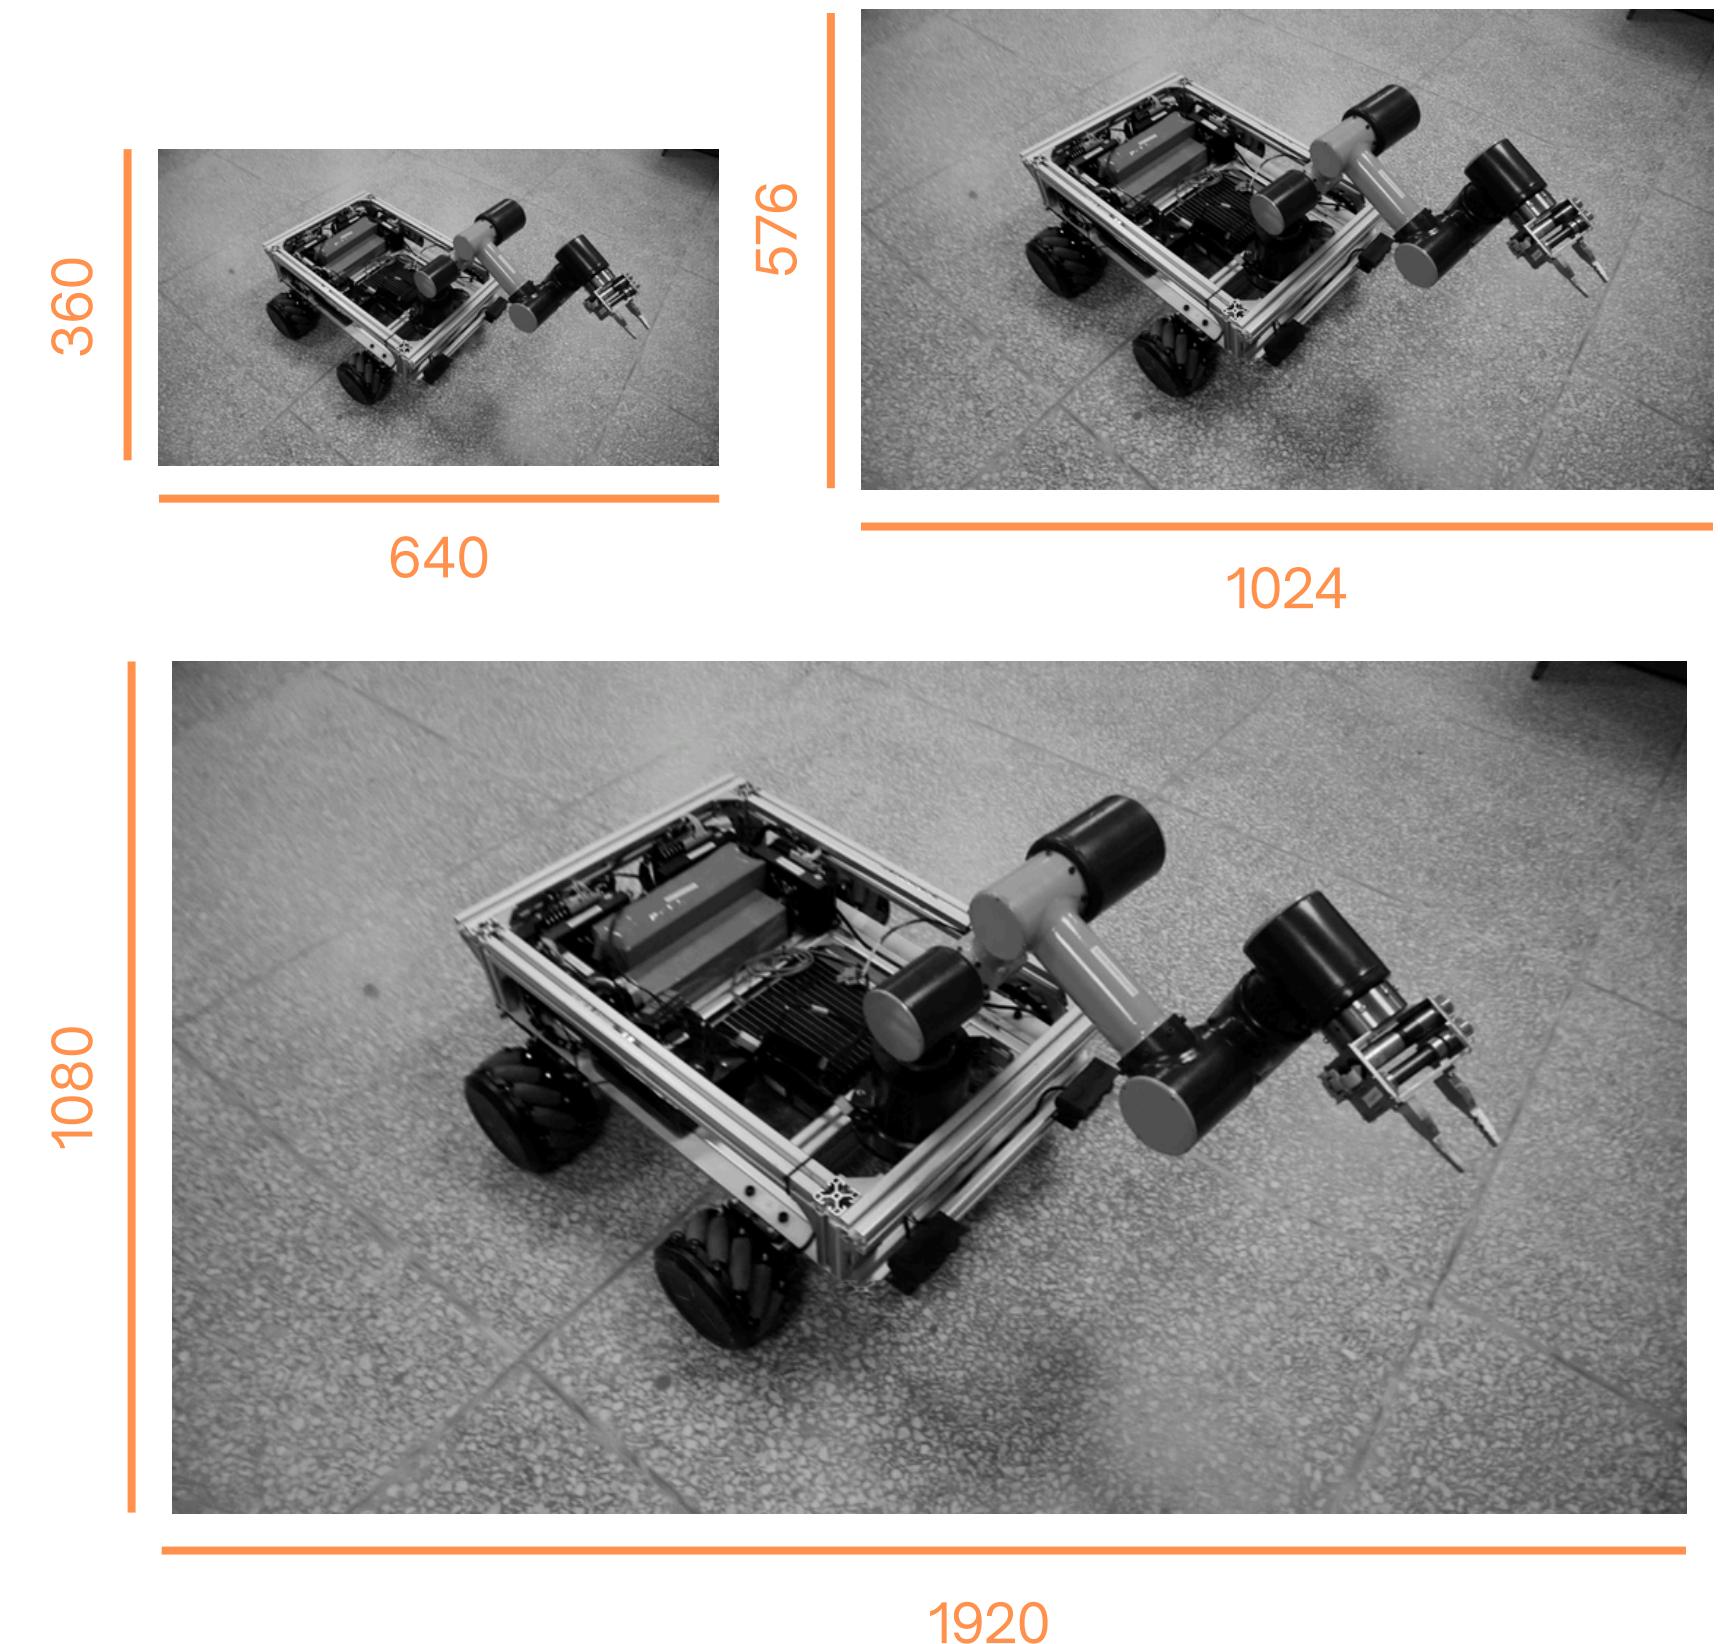
3 workloads:

- **Small:** 640x360 px
- **Medium:** 1024x576 px
- **Large:** 1920x1080 px

30 frames each

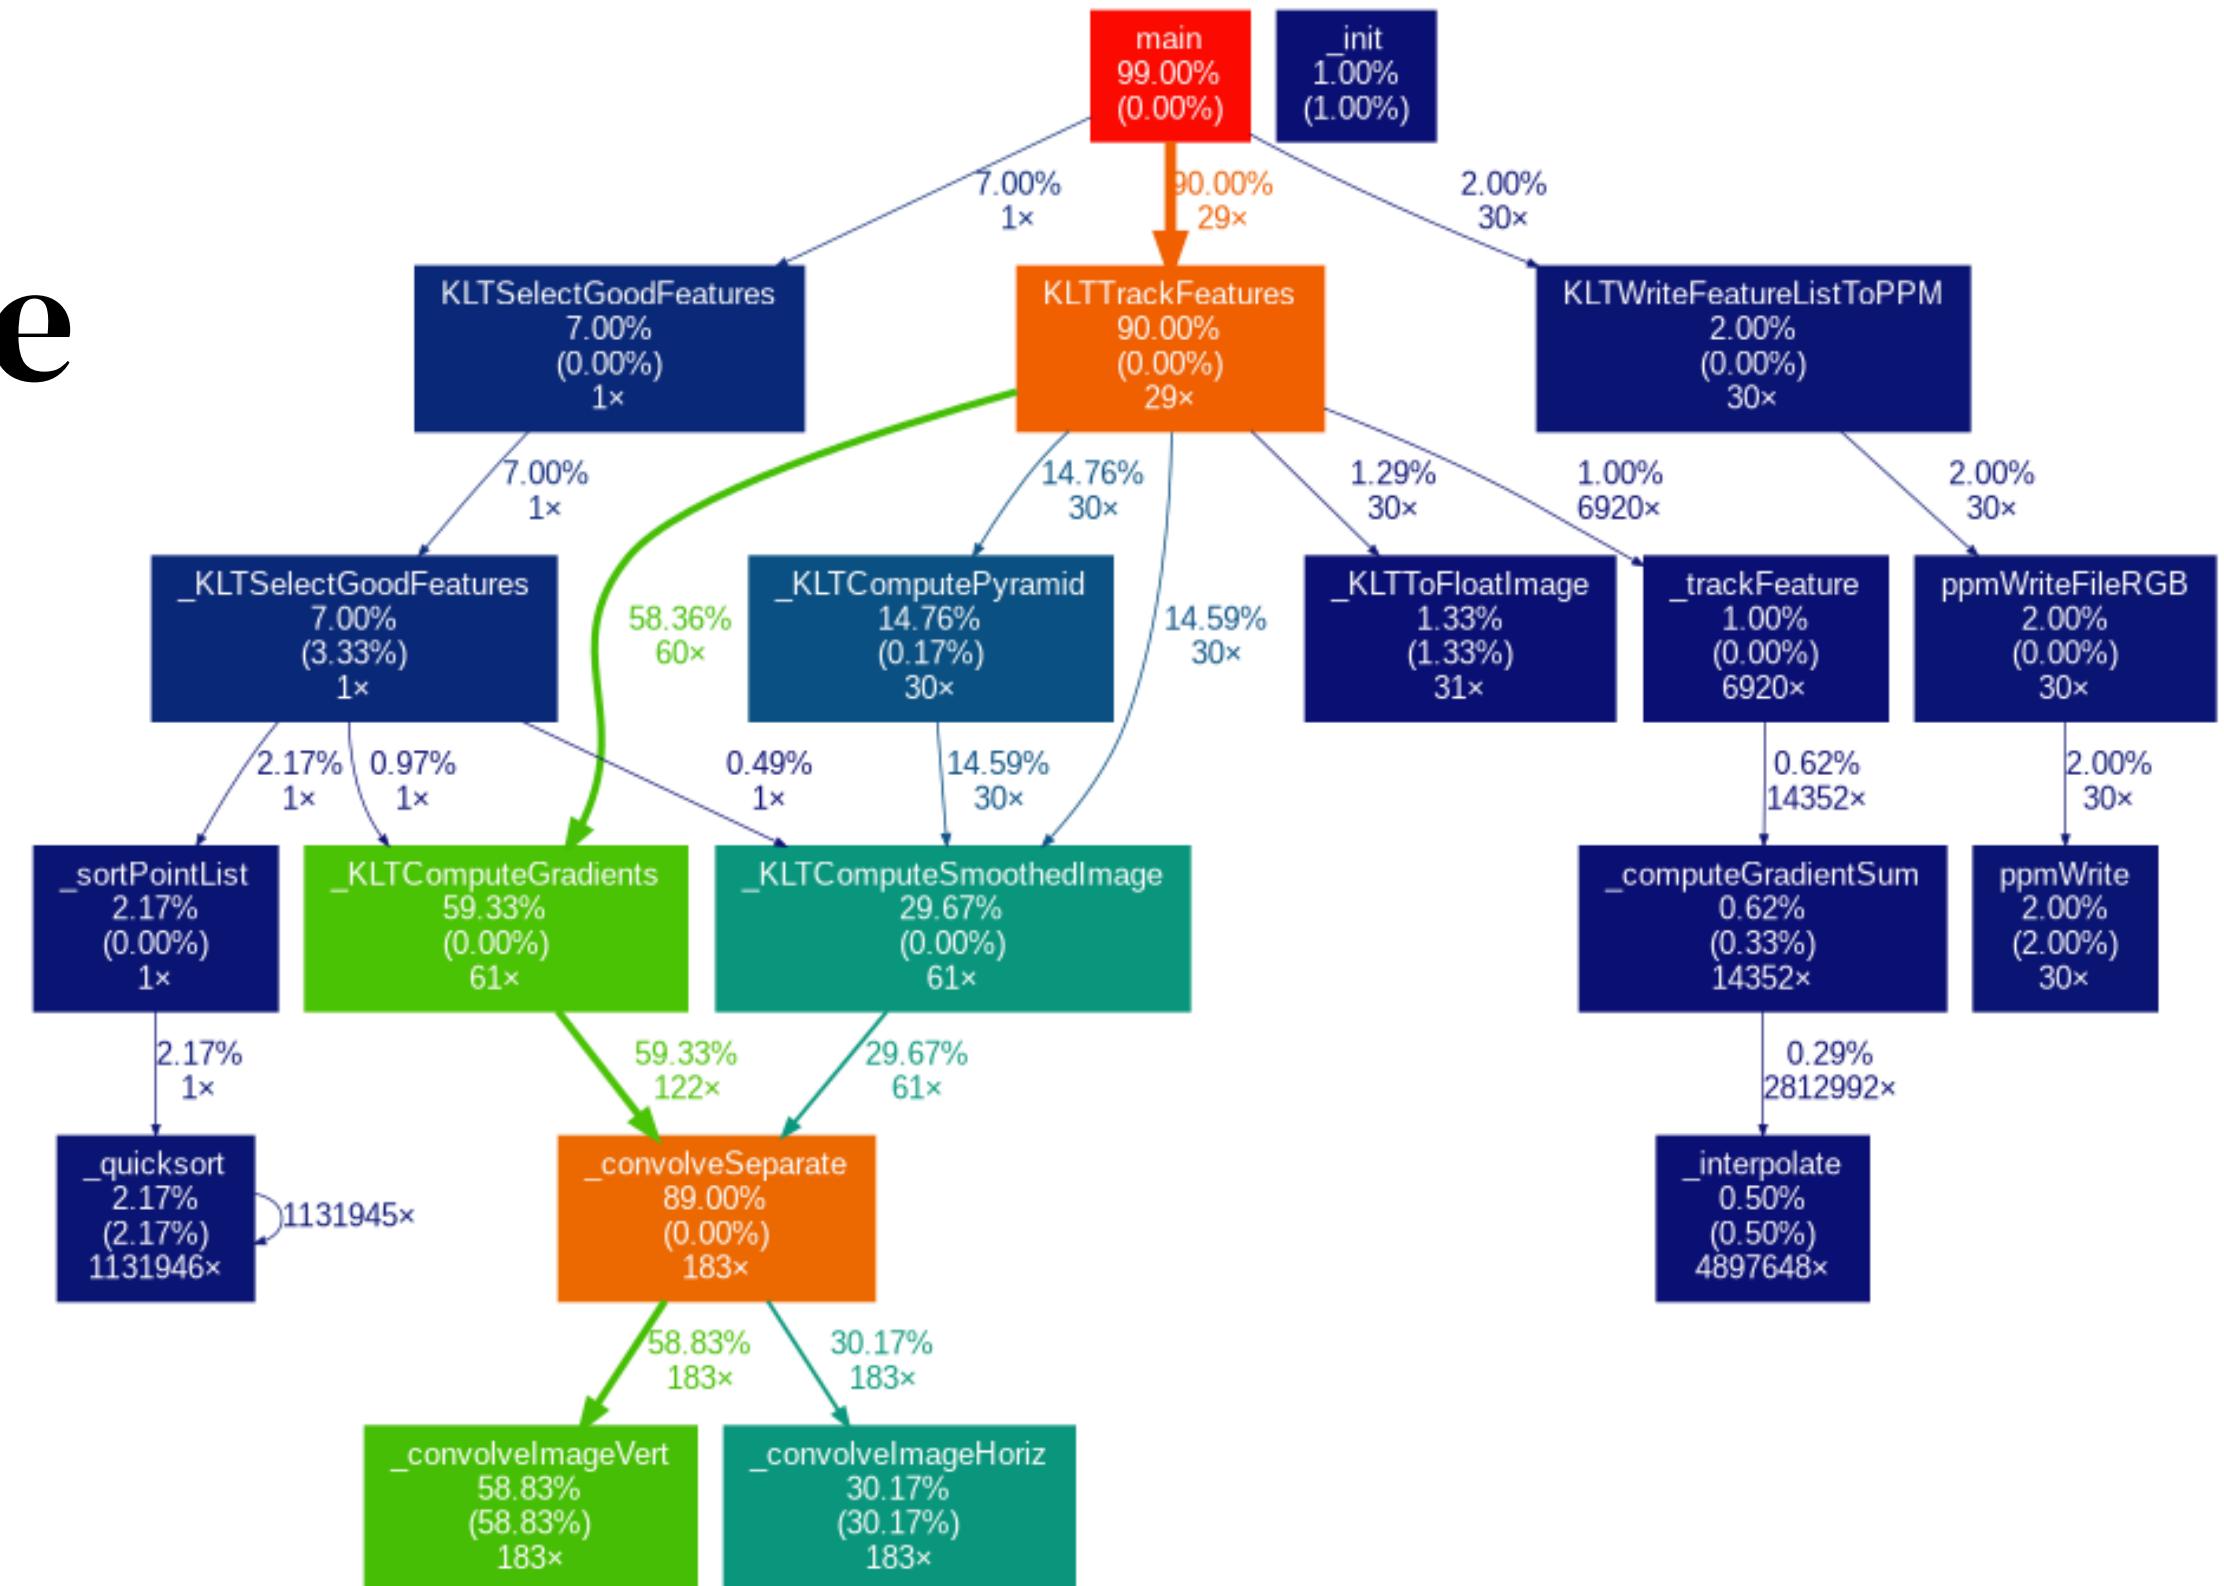
Running on:

NVIDIA GeForce RTX 3080 (Compute Capability 8.6)



Motivation

V1: CPU baseline



gprof2dot results on V1

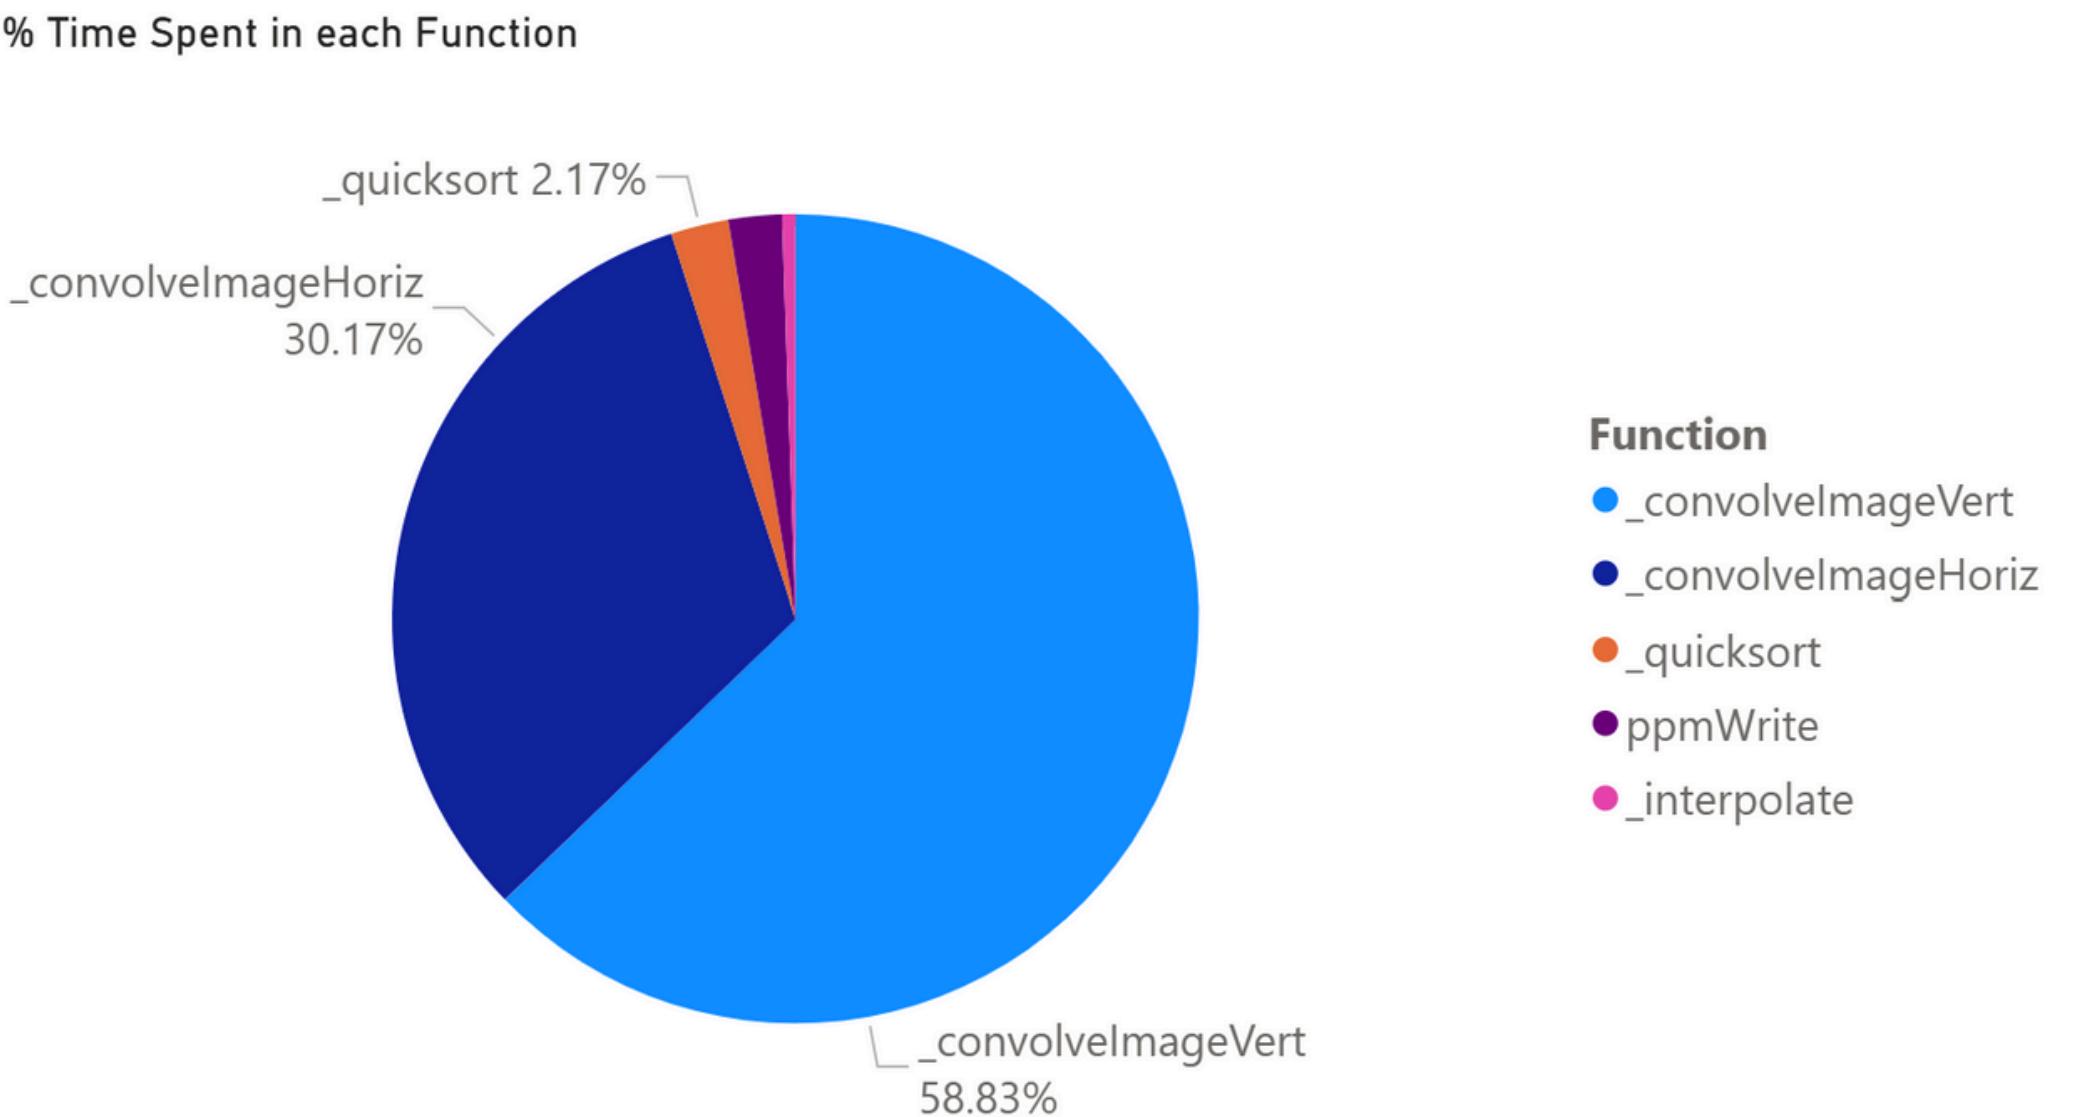
Our bottlenecks...

We ran our program several times over different datasets to find our main bottlenecks:

- `_convolveImageHoriz`
- `_convolveImageVert`
- `_interpolate` *

* `_interpolate` is included because it gets called many times

Identified **Convolutions** and **Feature Tracking** as main opportunities for parallel computation



V2: Naive GPU Port

Introduction of kernels:

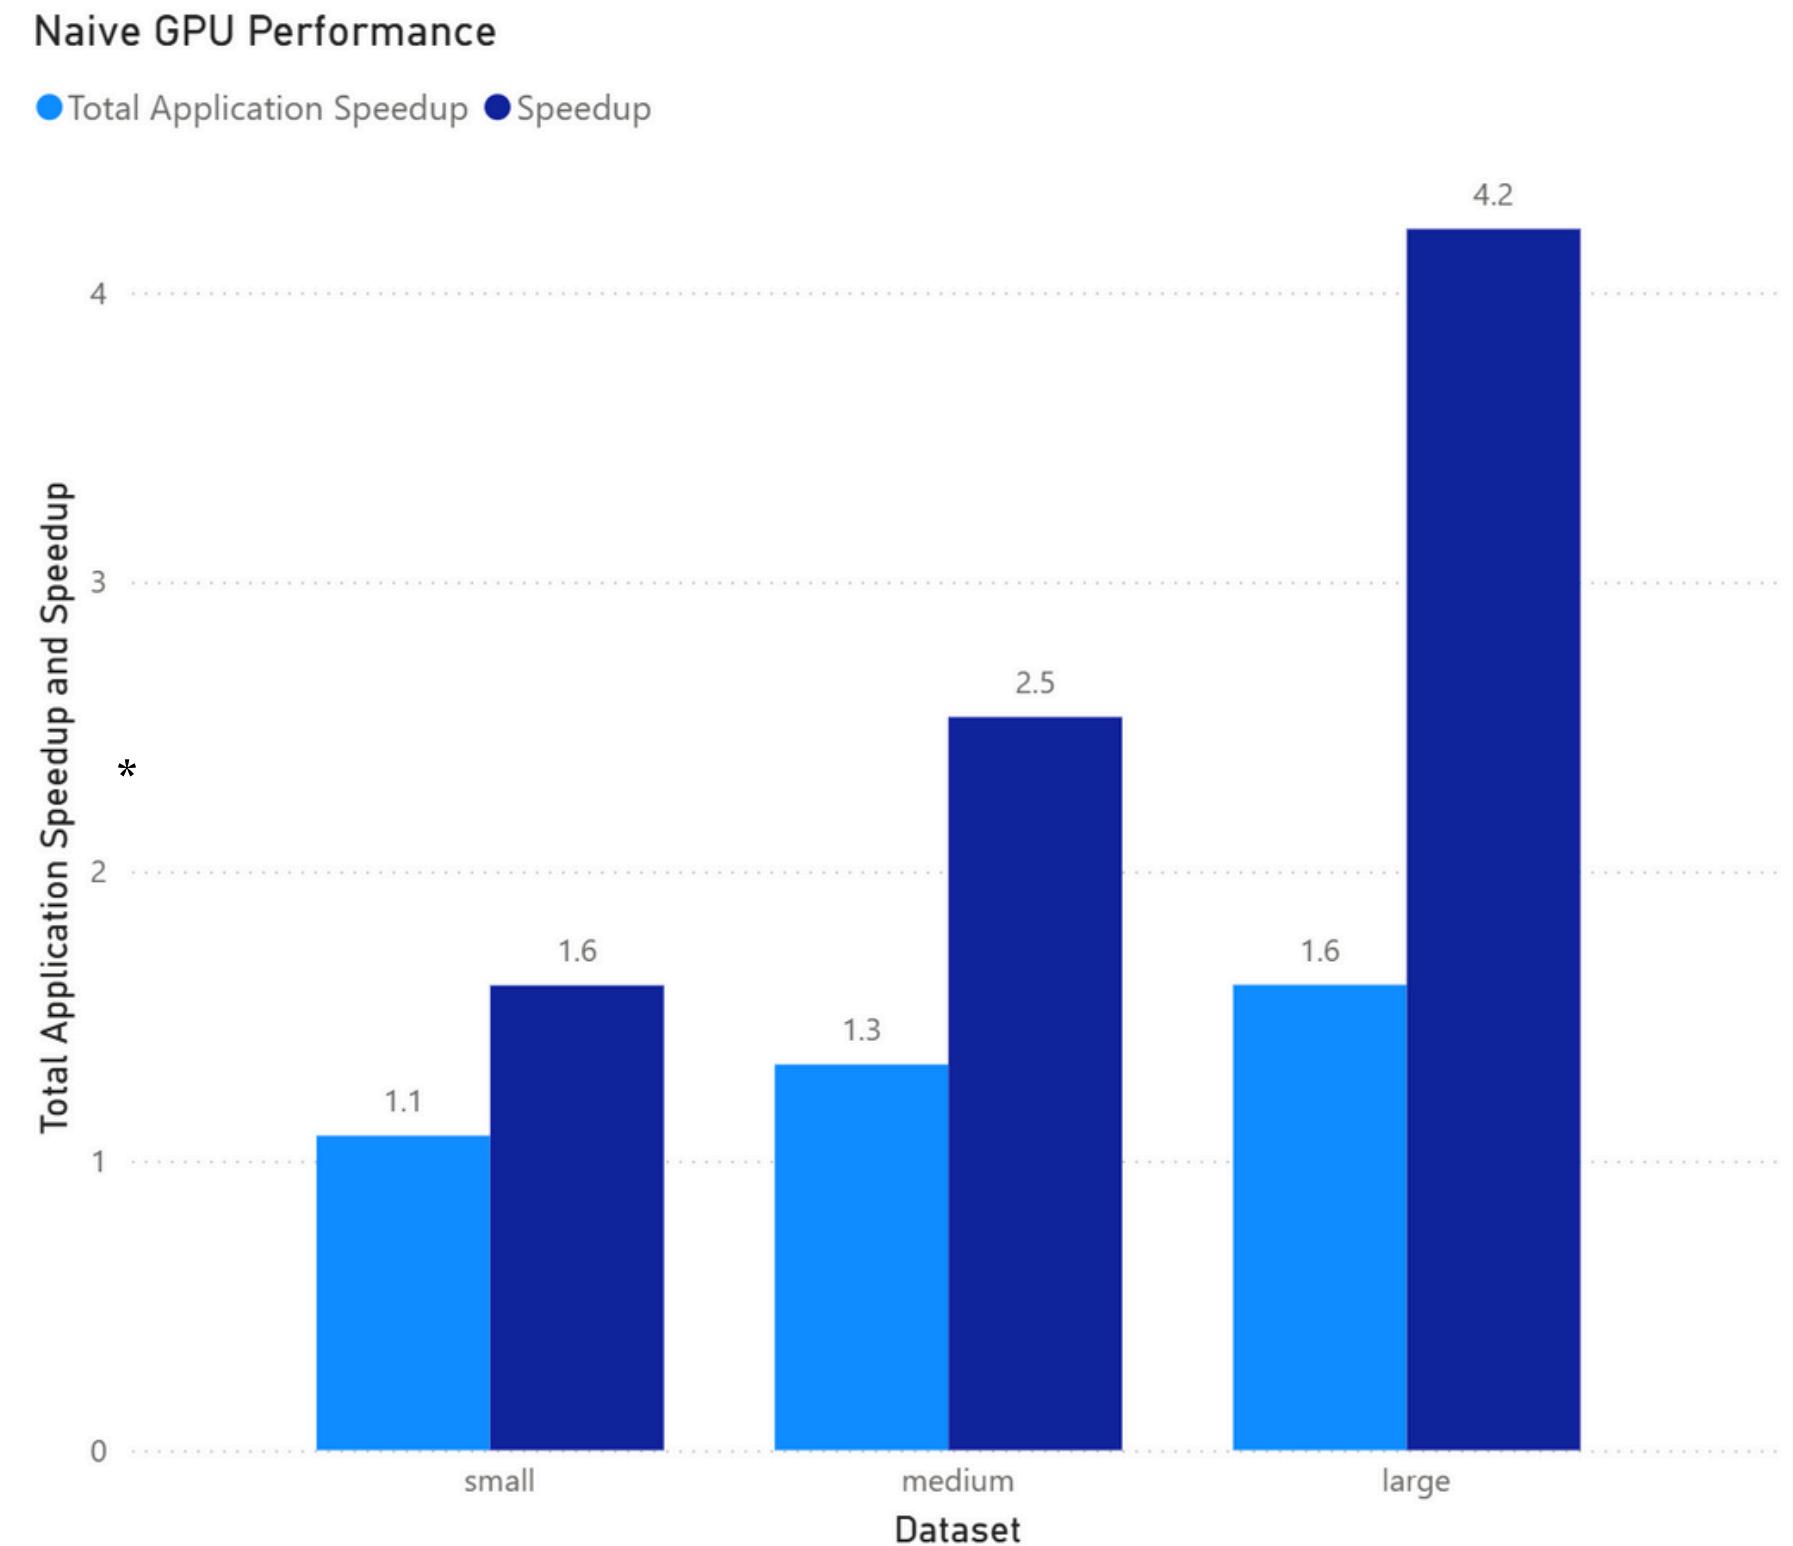
- convolveHorizKernel
- convolveVertKernel
- trackFeatureKernel

Bottlenecks:

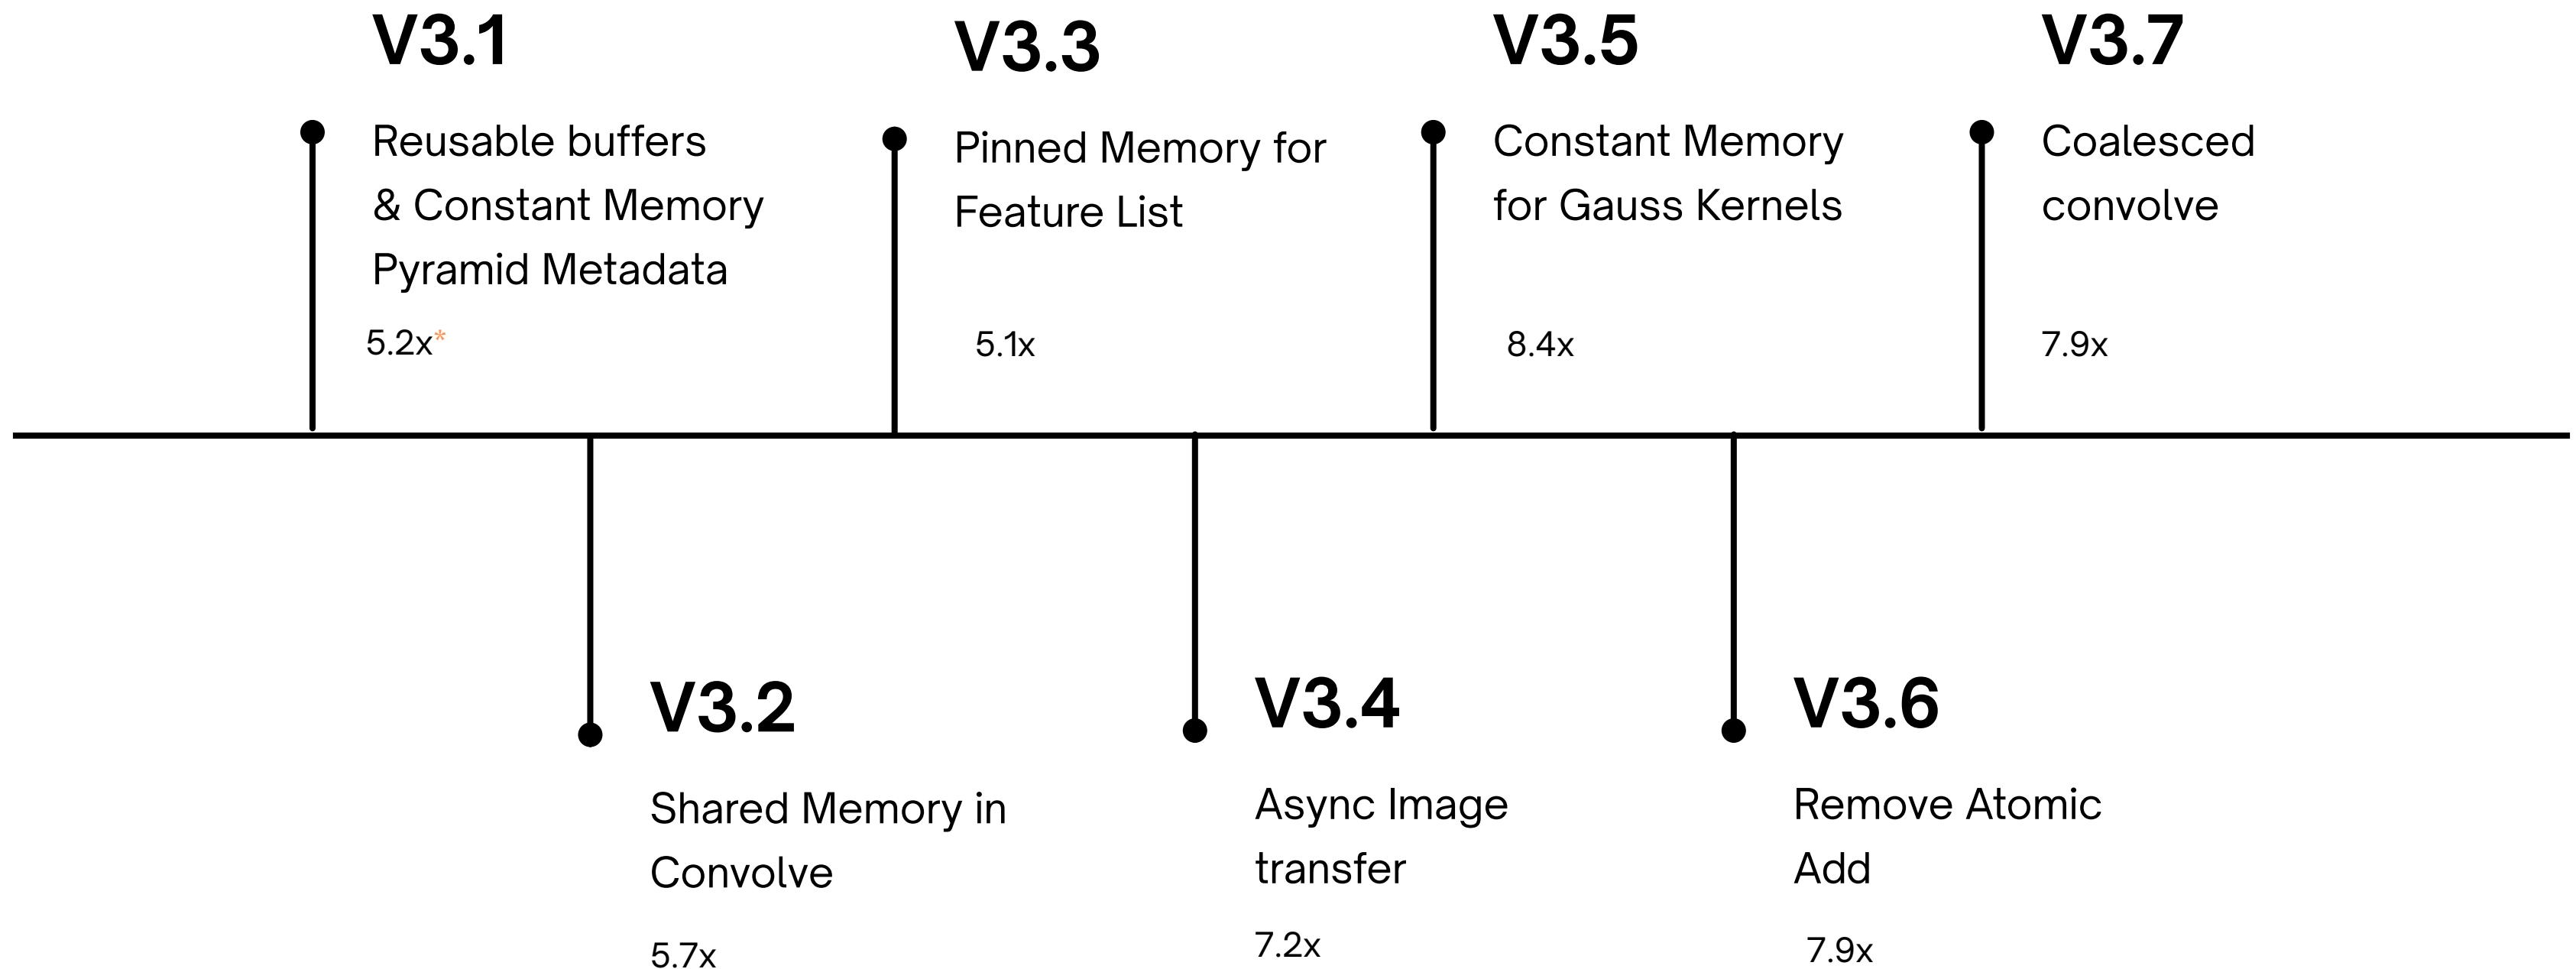
- Memory transfers (H2D, D2H) took upto **51.2%** of tracking time (large dataset)

Small Gains:

- **1.6x** speedup for total application time
- **4.2x** speedup if only tracking time is considered



V3: Optimising V2...



*on our medium workload

V3: Optimised!

Obtained up to **18.6x** speedup!

Small:

- **2.5x** tracking time
- **1.4x** total application time

Medium:

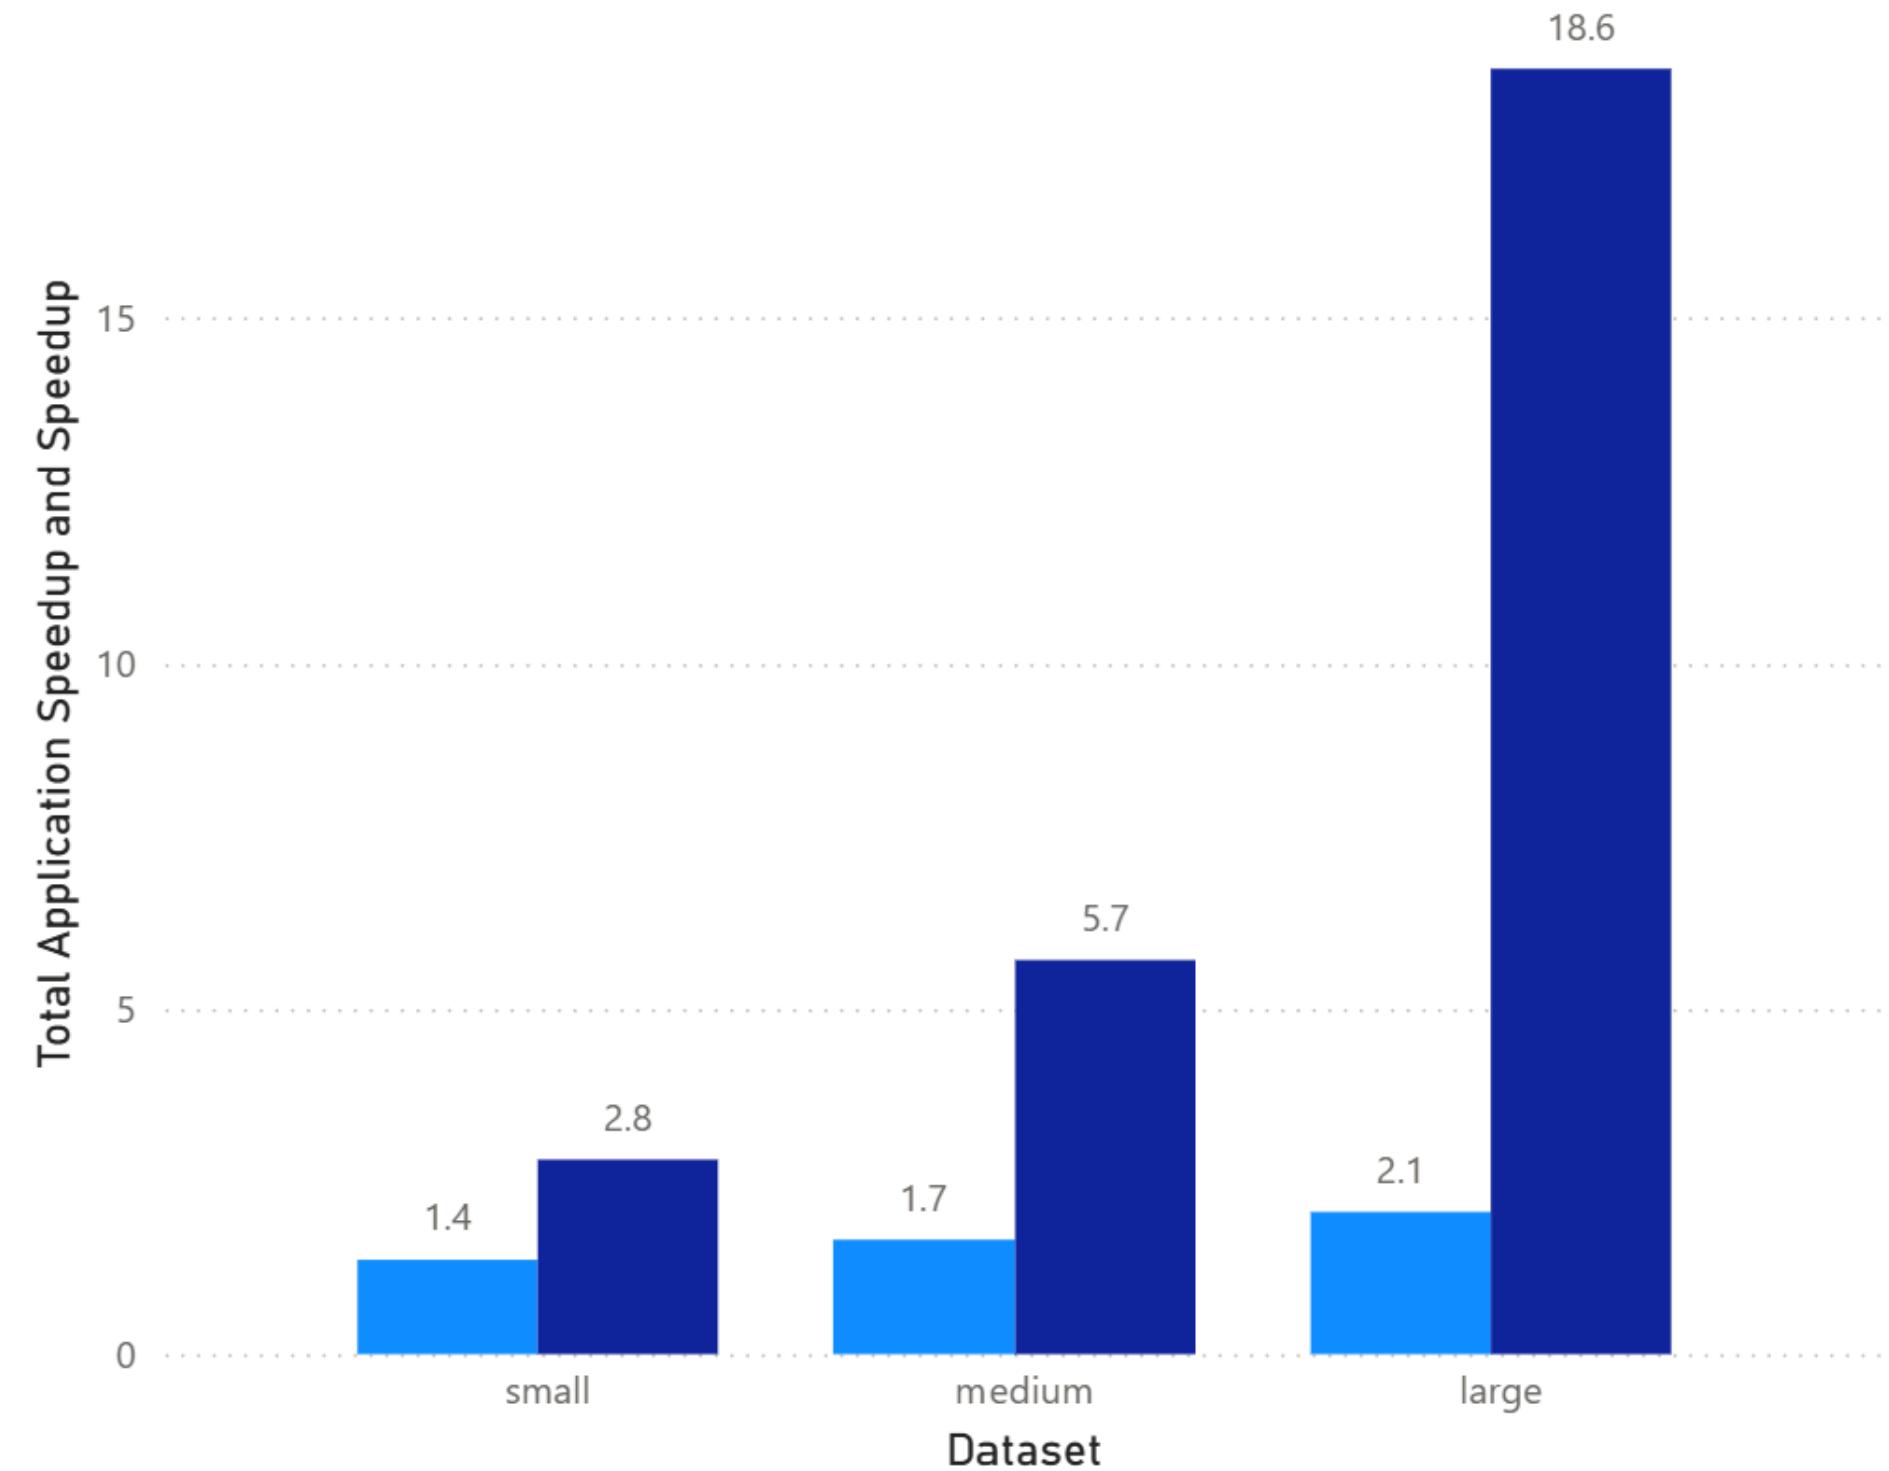
- **5.7x** tracking time
- **1.7x** total application time

Large

- **18.6x** tracking time
- **2.1x** total application time

Optimised GPU Performance

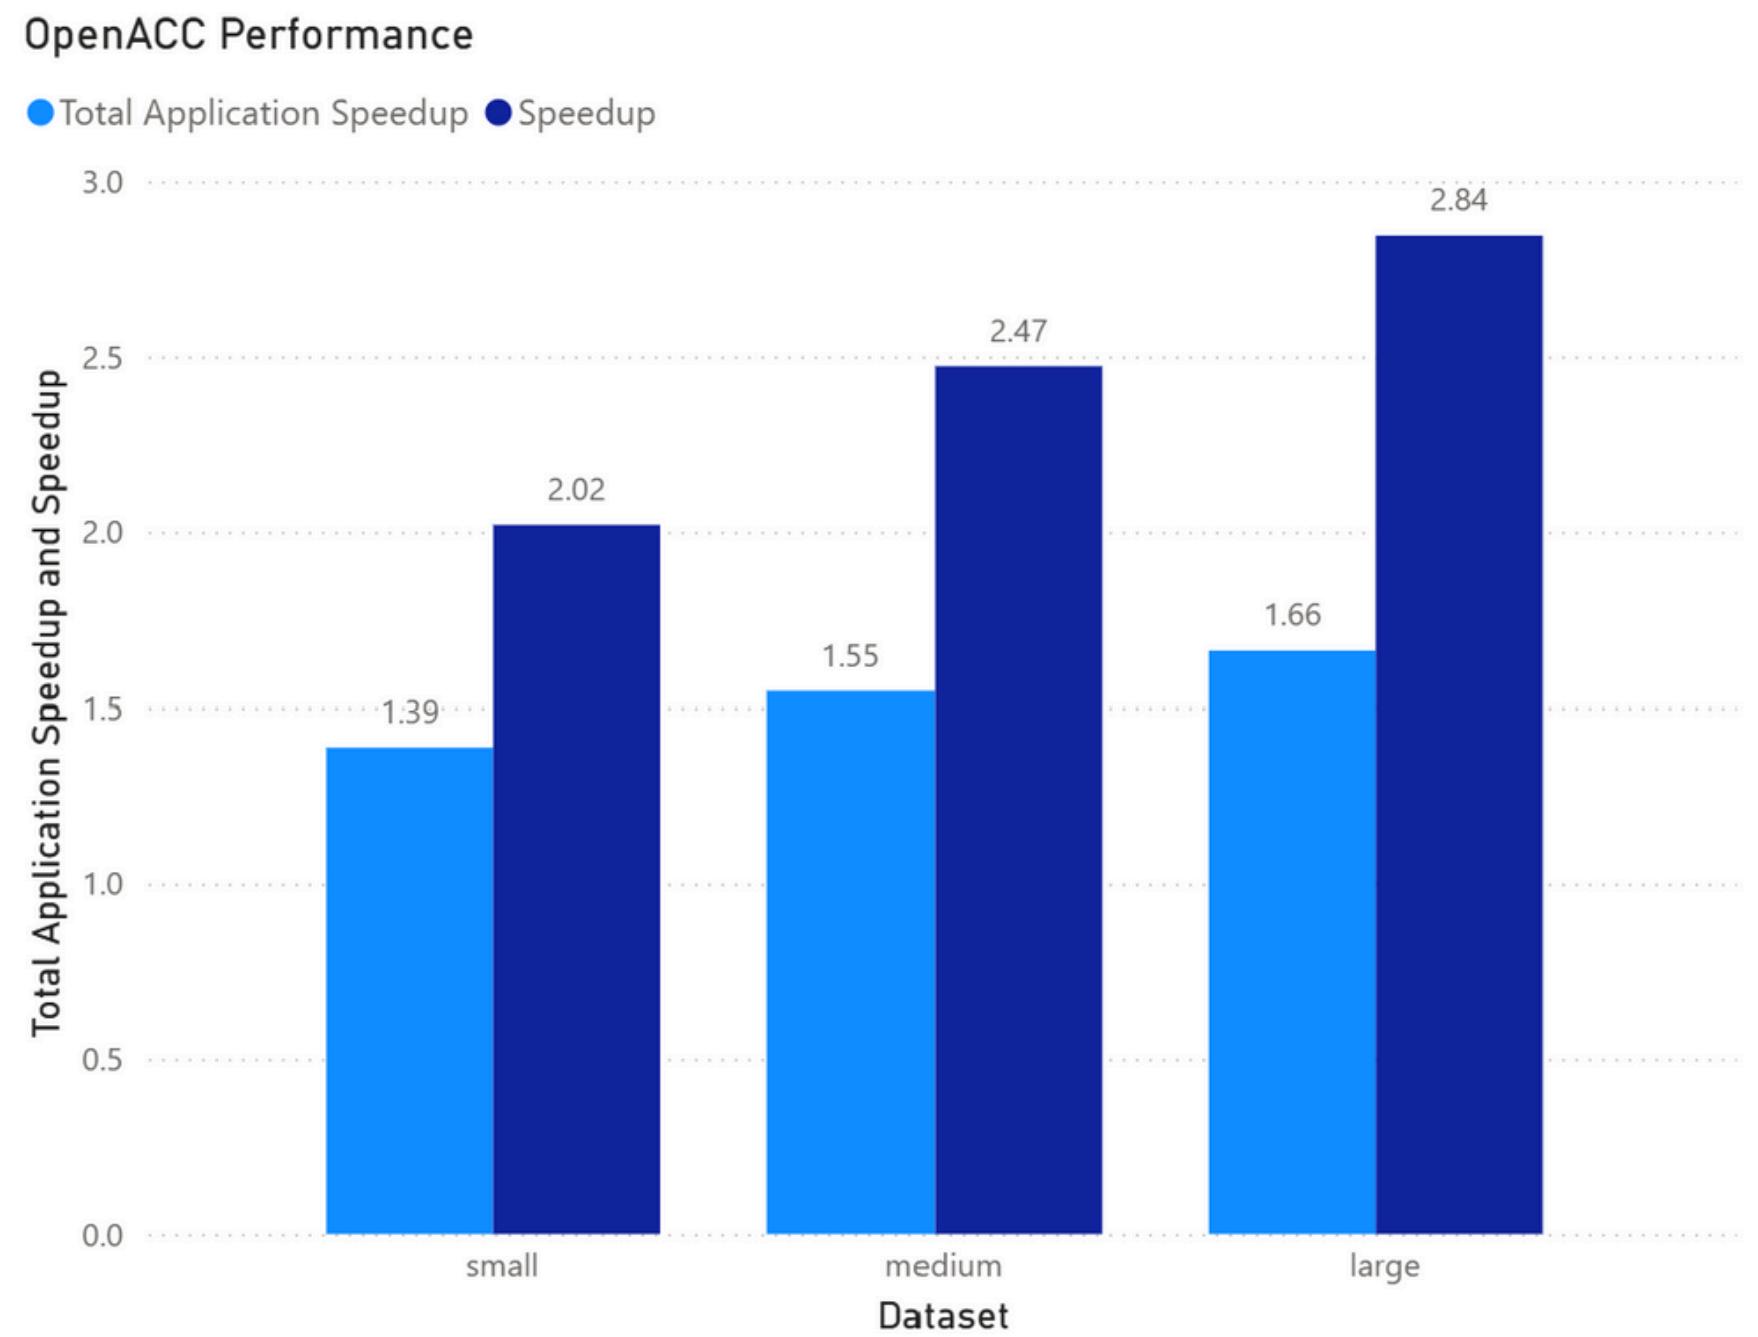
● Total Application Speedup ● Speedup



V4: OpenACC

Paralleling our main bottlenecks:
Horizontal and Vertical Convolution
using pragma ACC kernels

Quick, but Small Gains: 2.84x
on large dataset
(1.66x application speedup)

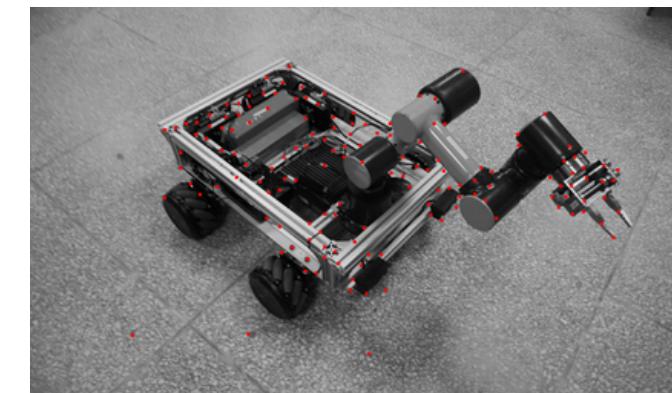


Correctness

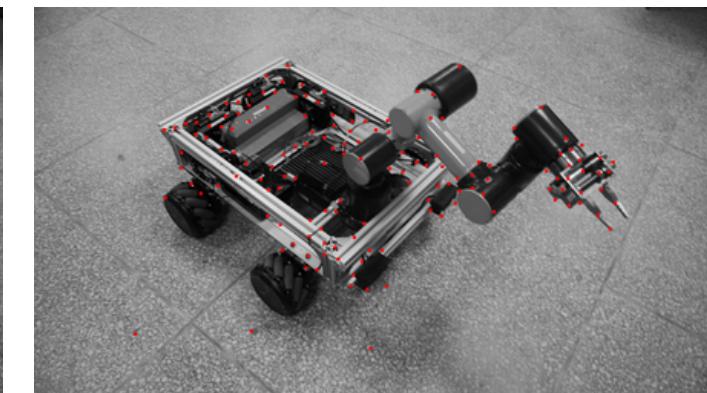
CPU tracks:

- small : 150 features
- medium : 145 features
- large: 113 features

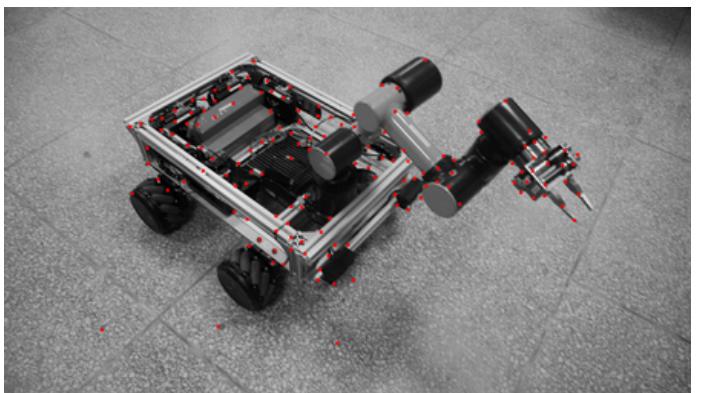
While V2 shows slight divergence, V3 and V4 track features correctly when compared with CPU



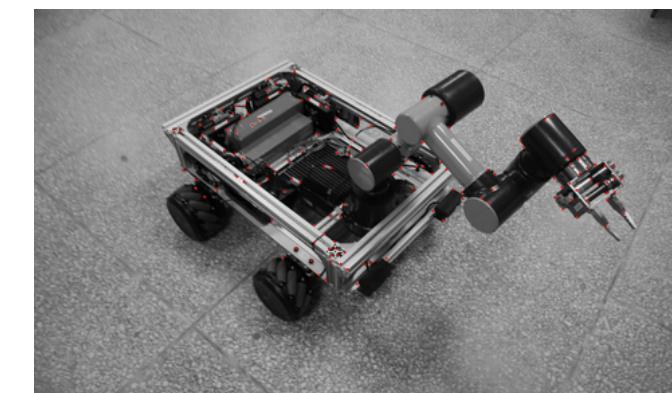
frame 0



frame 14



frame 29



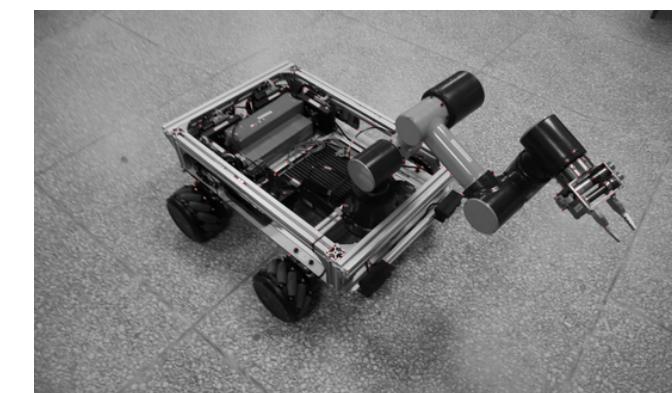
frame 0



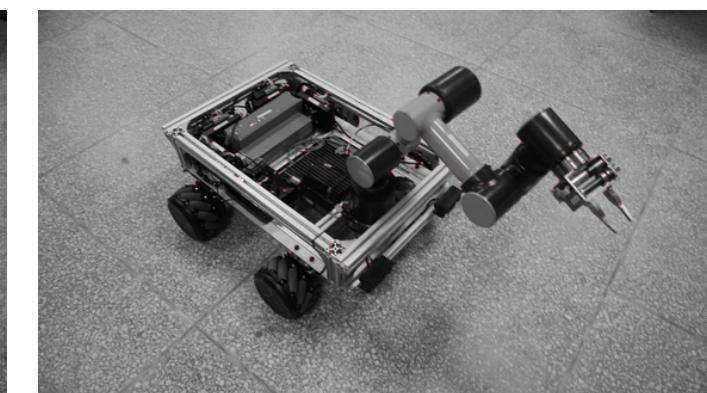
frame 14



frame 29



frame 0



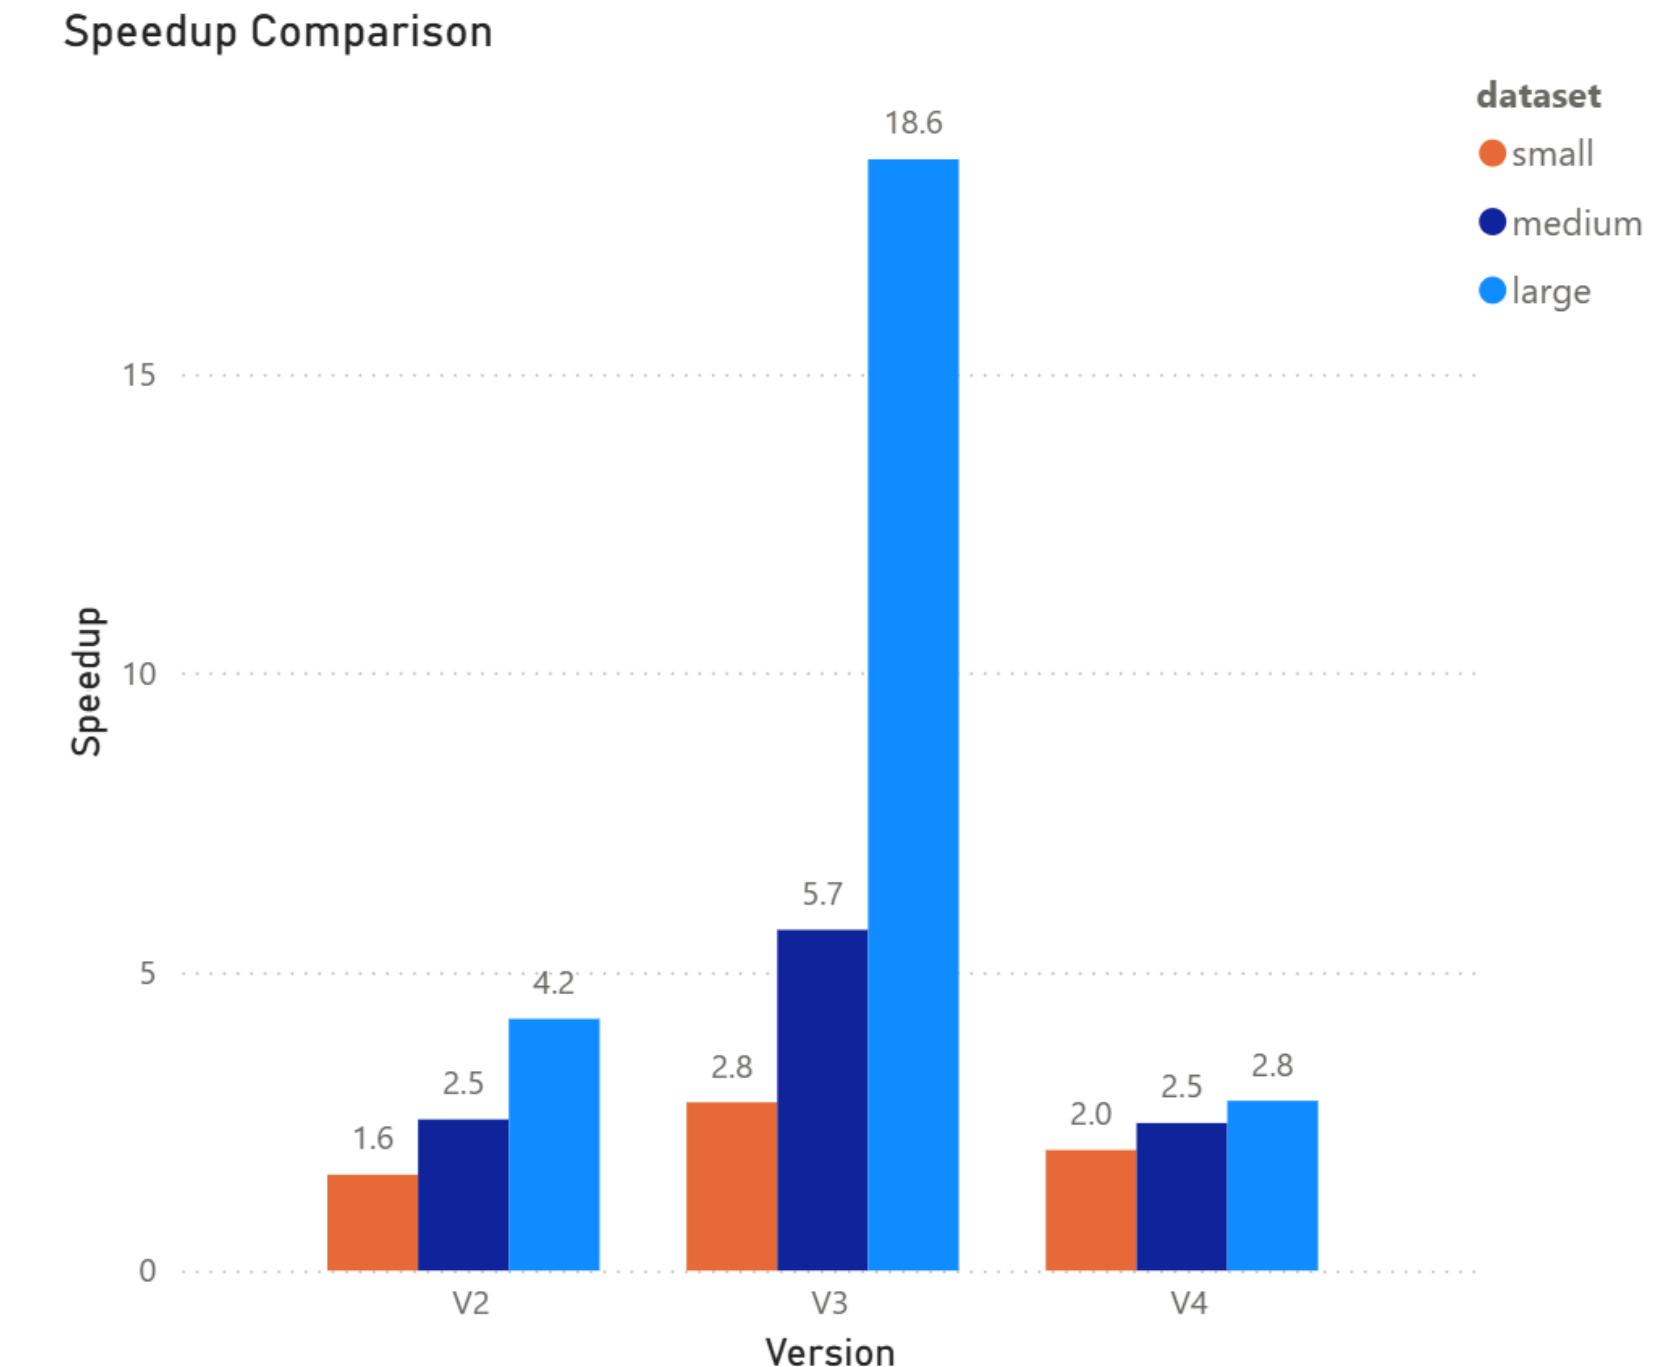
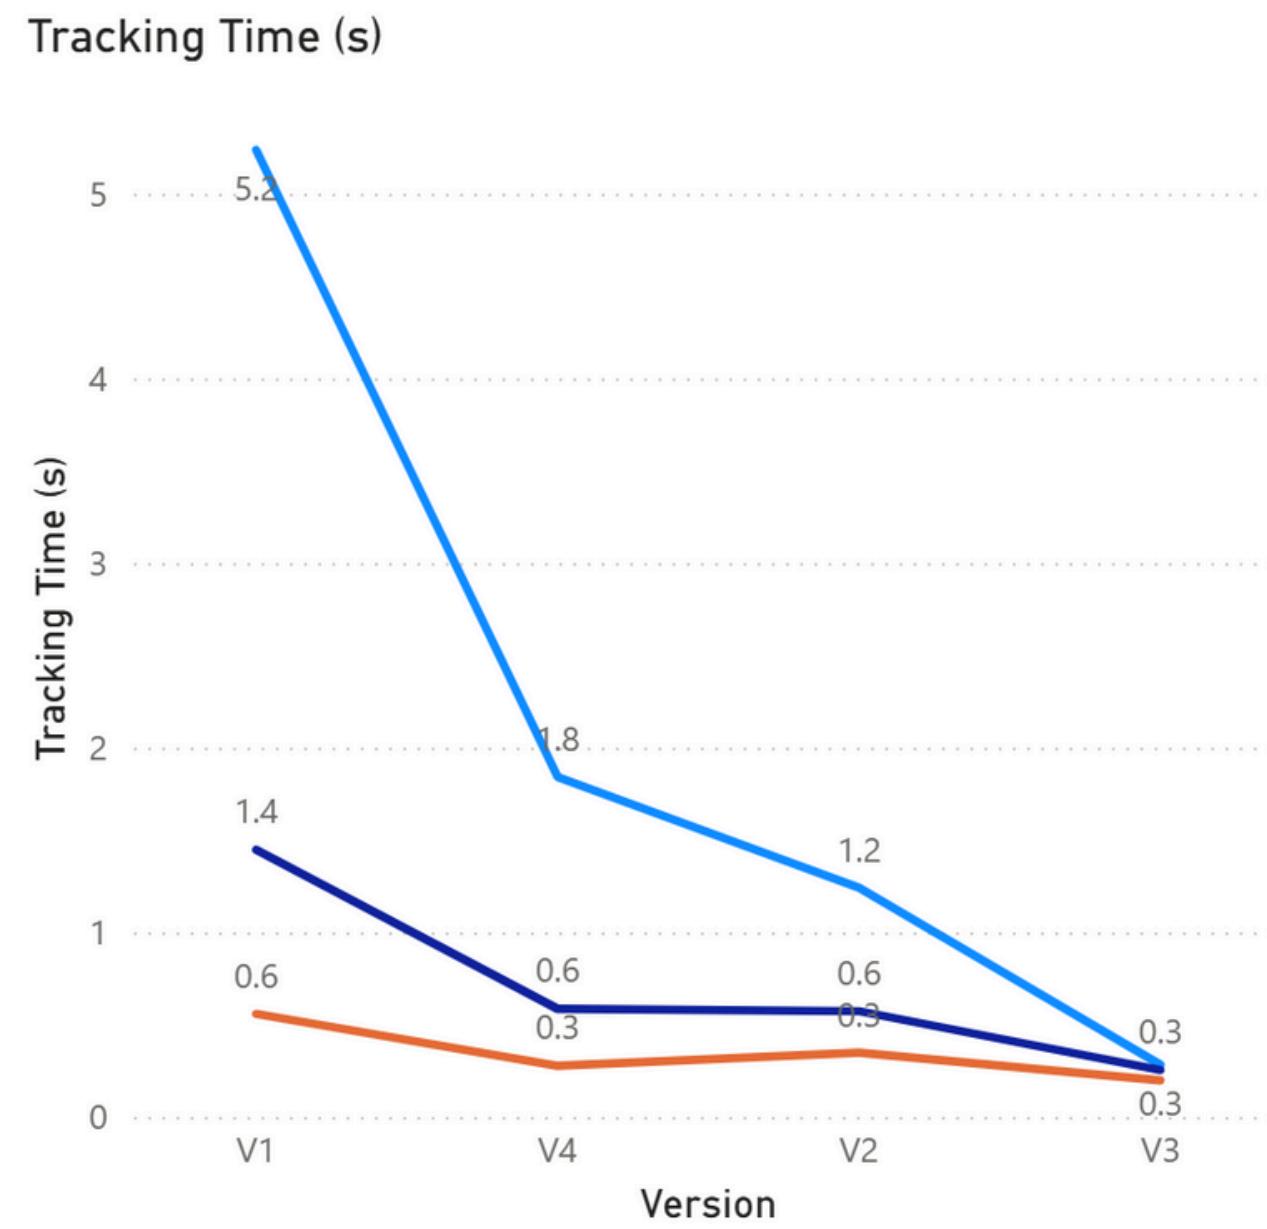
frame 14



frame 29

Comparison of Results

Speedup



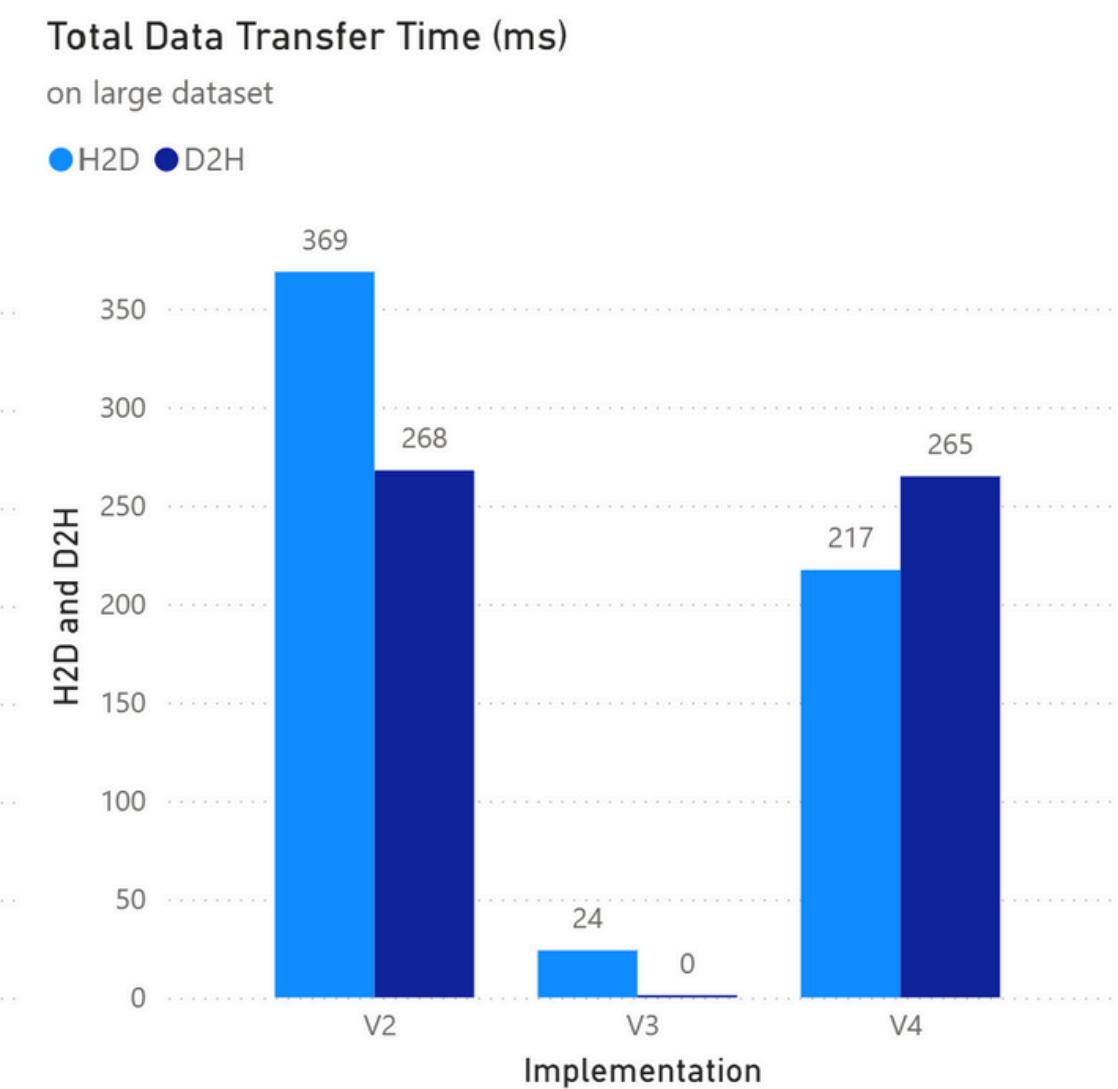
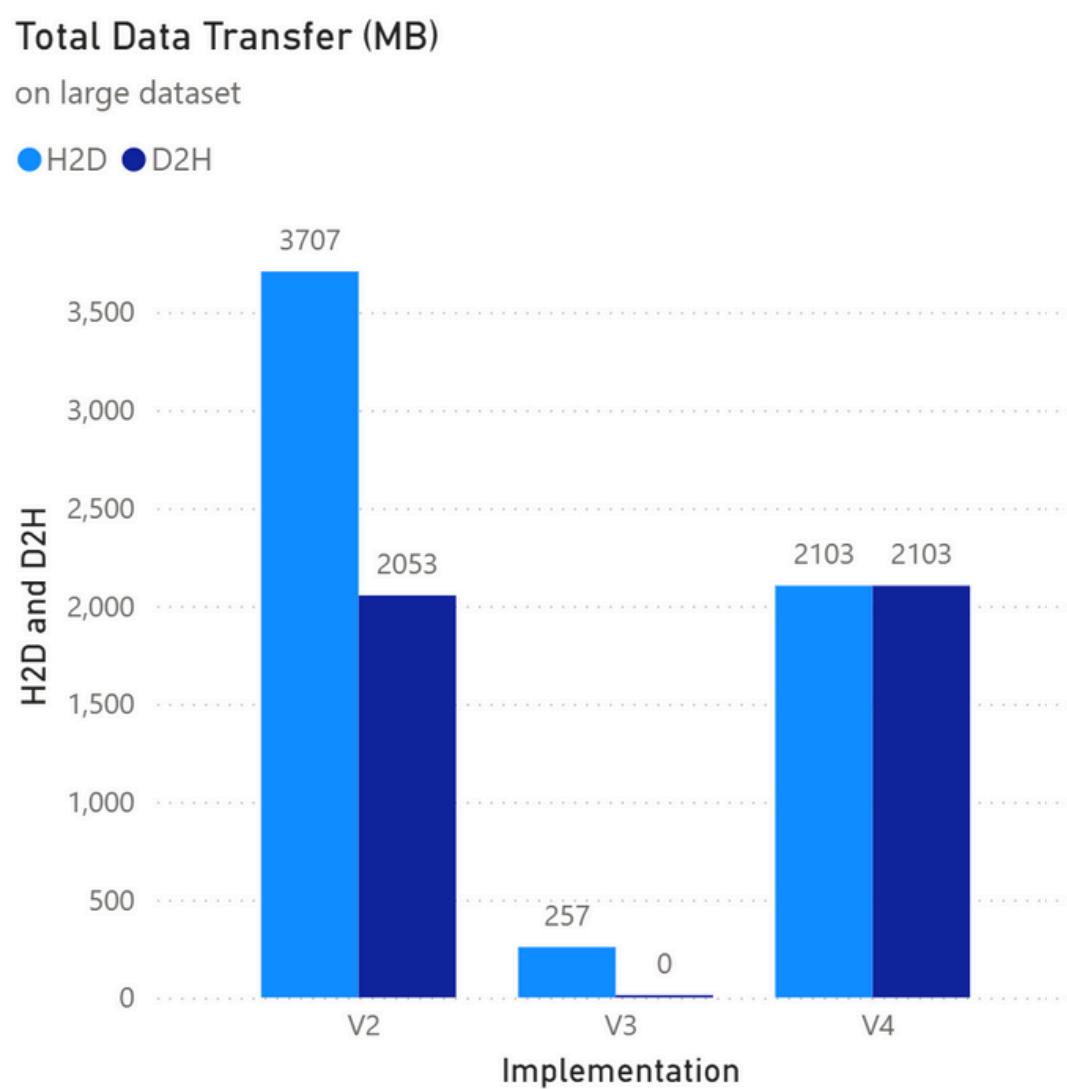
CUDA Brought Tracking Time down from **5.2s** to **0.28s**
while **OpenACC** brought it down to **1.84s**. *

OpenACC shows poor scaling
* on large dataset

Comparison of Results

Memory Transfer

The amount of data (MBs), and the time taken transferring that data (ms) from Host to Device and Device to Host **decrease** significantly from our CUDA optimization



Comparison of Results

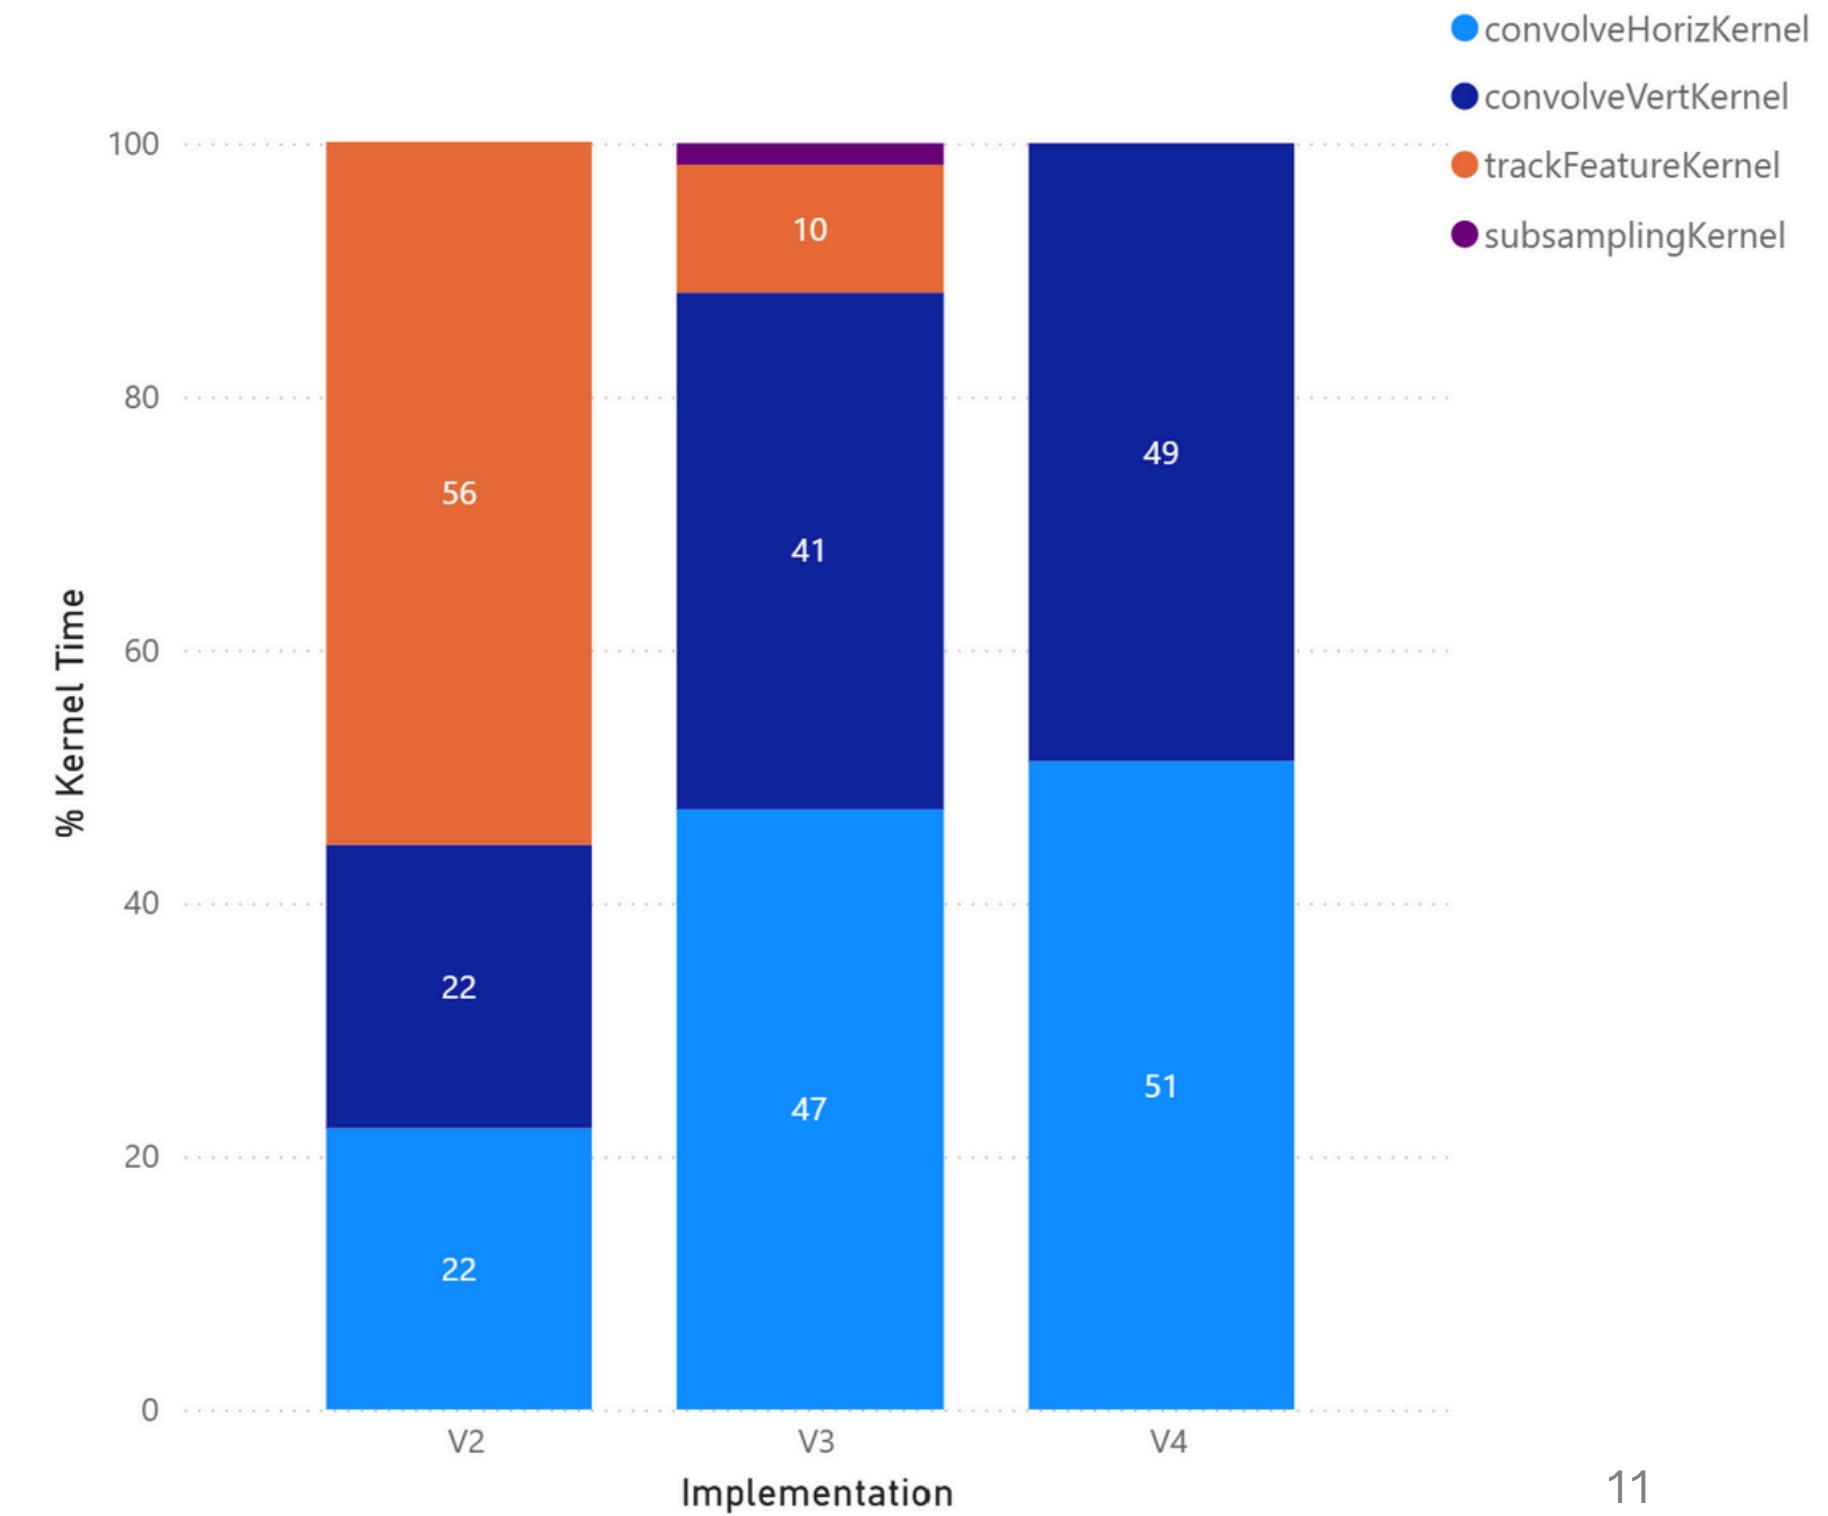
Kernel Breakdown

Before Optimisation:
Tracking Dominates

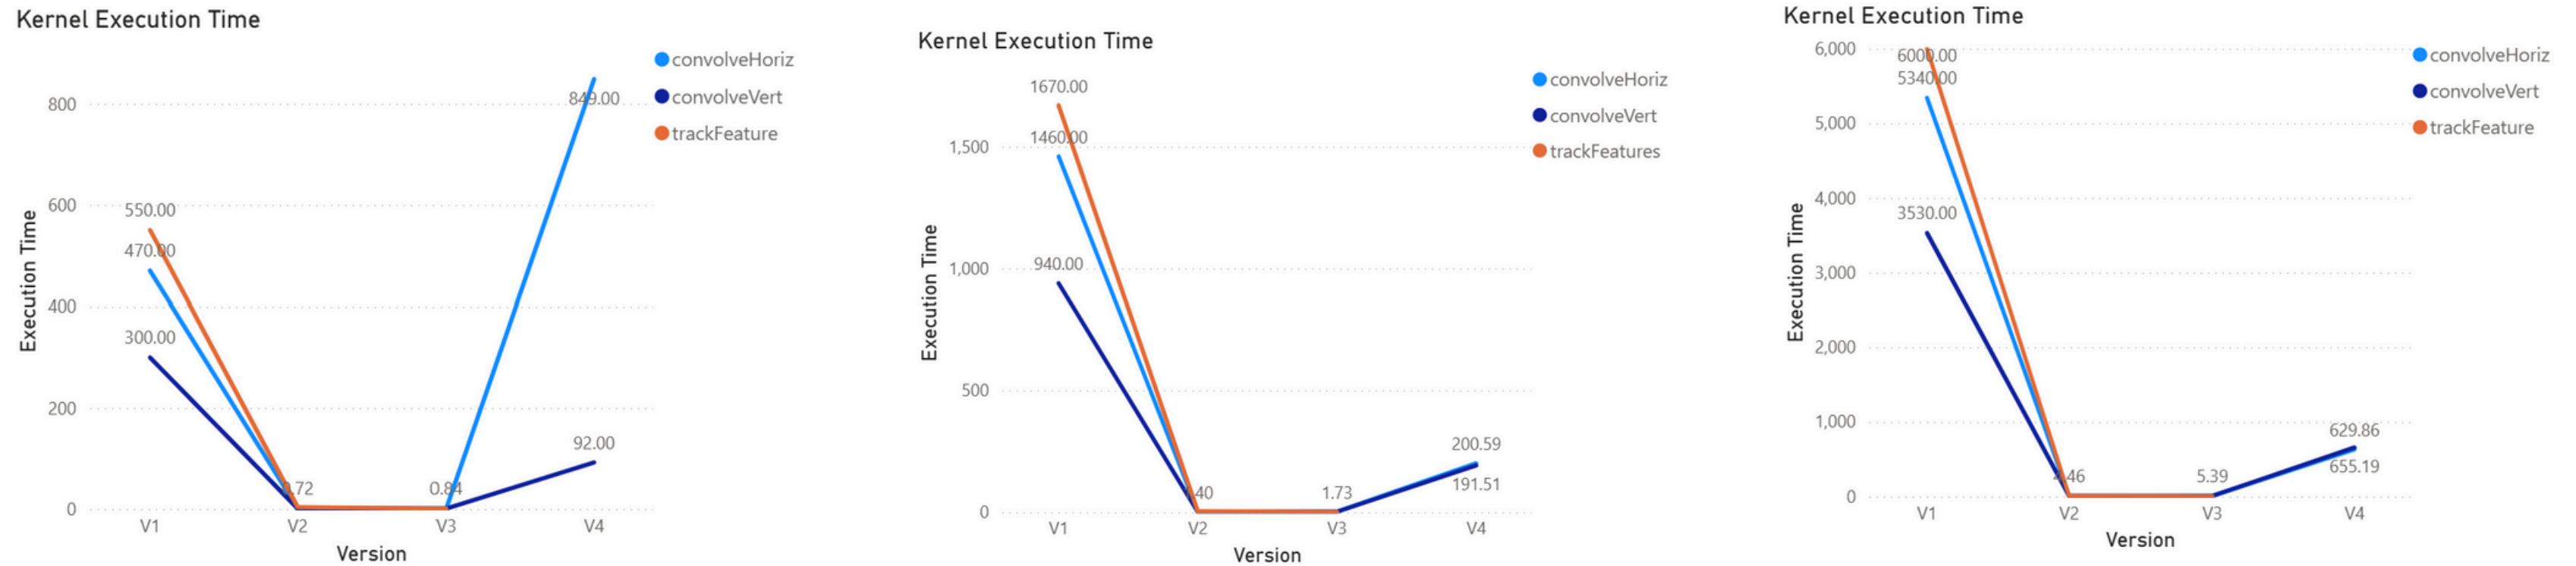
After Optimisation:
Convolutions Dominate

This aligns with the CPU baseline!

Kernel Breakdown Summary
on medium dataset



Comparison of Results



Successfully parallelised our bottlenecks:

Horizontal Convolution: **5340 ms** → **5.39 ms**

Vertical Convolution: **3530 ms** → **4.93 ms**

Feature Tracking: **6000 ms** → **0.39 ms**

Key Takeaways

- Highly Suitable for Parallelisation
- Scales well for large workloads
- CUDA > OpenACC?
- Sequential Dependencies Constraints